

Bachelor's Degree Program in Software Engineering



Prof. Dr. Paulo Ditarso Maciel Jr. (paulo.maciел@ifpb.edu.br)

September 1, 2026

Operating Systems

Based on material from the book *Operating Systems: Three Easy Pieces*.

```
$ git clone https://github.com/pdmjr/ostep\_slides.git
```

Disclaimer¹

- This presentation is not intended to dive deeply into any topic.
- It is essential to read the books and materials provided!
- What is the purpose of this presentation?



Agenda

1. A Dialogue on the Book
2. Introduction
3. A Dialogue on Virtualization
4. Processes
5. Process API
6. Direct Execution
7. CPU Scheduling

A Dialogue on the Book

A Dialogue on the Book

The dialogue introduces the book and explains the origin of its title, inspired by Richard Feynman's *Six Easy Pieces*.

The book is organized around three fundamental ideas: **virtualization**, **concurrency**, and **persistence**. Through these topics, it explores how operating systems manage processors, memory, storage, virtual machines, and even some aspects of distributed systems.

The dialogue also emphasizes an effective approach to learning operating systems: attending classes, reviewing the material regularly, completing homework, and especially developing projects that involve writing real code to solve real problems.

Finally, the dialogue explains that these conversations are intended to encourage reflection on complex concepts and help the reader reason about the ideas presented throughout the book.

Introduction

What Happens When a Program Runs?

- A running program executes instructions.
- The processor repeatedly **fetches**, **decodes**, and **executes** instructions.
- Programs assume an orderly, sequential execution model.
- The operating system works underneath this simple model to make the machine easier to use.

Key idea

The OS hides much of the complexity involved in sharing and managing hardware resources.

What Does an Operating System Do?

- Makes it easier to run programs and use hardware resources.
- **Virtualizes** physical resources such as the CPU, memory, and storage.
- Provides interfaces through **system calls**.
- Acts as a **resource manager** for shared system resources.

The Crux of the Problem

How does the operating system virtualize resources efficiently, and what hardware support is needed?

Virtualizing the CPU

- The example program repeatedly waits and prints a user-provided string.
- A single program running alone is straightforward.
- The interesting behavior appears when several copies execute at once.

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <sys/time.h>
4  #include <assert.h>
5  #include "common.h"
6
7  int
8  main(int argc, char *argv[])
9  {
10     if (argc != 2) {
11         fprintf(stderr, "usage: cpu <string>\n");
12         exit(1);
13     }
14     char *str = argv[1];
15     while (1) {
16         Spin(1);
17         printf("%s\n", str);
18     }
19     return 0;
20 }
```

Figure 2.1: Simple Example: Code That Loops And Prints (`cpu.c`)

Source: chapter page 3.

The Illusion of Many CPUs

```
prompt> ./cpu A & ./cpu B & ./cpu C & ./cpu D &  
[1] 7353  
[2] 7354  
[3] 7355  
[4] 7356  
A  
B  
D  
C  
A  
B  
D  
C  
A  
...
```

Figure 2.2: Running Many Programs At Once

Source: chapter page 4.

- Multiple programs appear to run simultaneously even on one processor.
- The OS, with hardware support, creates the illusion of many **virtual CPUs**.
- The OS also needs policies to decide which program should run at a given time.

Virtualizing Memory

- Memory can be viewed as an array of bytes addressed by the processor.
- Programs continuously access memory for instructions and data.
- The example allocates memory with *malloc()*, initializes it, and repeatedly updates it.

```
1  #include <unistd.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include "common.h"
5
6  int
7  main(int argc, char *argv[])
8  {
9      int *p = malloc(sizeof(int));           // a1
10     assert(p != NULL);
11     printf("(%d) address pointed to by p: %p\n",
12            getpid(), p);                     // a2
13     *p = 0;                                  // a3
14     while (1) {
15         Spin(1);
16         *p = *p + 1;
17         printf("(%d) p: %d\n", getpid(), *p); // a4
18     }
19     return 0;
20 }
```

Figure 2.3: A Program That Accesses Memory (*mem.c*)

Source: chapter page 5.

Private Virtual Address Spaces

```
prompt> ./mem & ./mem &
[1] 24113
[2] 24114
(24113) address pointed to by p: 0x200000
(24114) address pointed to by p: 0x200000
(24113) p: 1
(24114) p: 1
(24114) p: 2
(24113) p: 2
(24113) p: 3
(24114) p: 3
(24113) p: 4
(24114) p: 4
...
```

Figure 2.4: Running The Memory Program Multiple Times

Source: chapter page 6.

- Two processes can appear to use the same address independently.
- Each process sees its own private **virtual address space**.
- The OS maps each virtual address space onto shared physical memory.
- One process normally does not affect another process's address space.

Concurrency

- Concurrency refers to problems that arise when multiple activities proceed at the same time.
- Modern programs may contain multiple threads sharing the same memory space.
- The example creates two worker threads that increment a shared counter.

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include "common.h"
4  #include "common_threads.h"
5
6  volatile int counter = 0;
7  int loops;
8
9  void *worker(void *arg) {
10     int i;
11     for (i = 0; i < loops; i++) {
12         counter++;
13     }
14     return NULL;
15 }
16
17 int main(int argc, char *argv[]) {
18     if (argc != 2) {
19         fprintf(stderr, "usage: threads <value>\n");
20         exit(1);
21     }
22     loops = atoi(argv[1]);
23     pthread_t p1, p2;
24     printf("Initial value : %d\n", counter);
25
26     Pthread_create(&p1, NULL, worker, NULL);
27     Pthread_create(&p2, NULL, worker, NULL);
28     Pthread_join(p1, NULL);
29     Pthread_join(p2, NULL);
30     printf("Final value   : %d\n", counter);
31     return 0;
32 }
```

Figure 2.5: A Multi-threaded Program (threads.c)

Source: chapter page 7.

Why Concurrency Is Difficult

- With 1,000 increments per thread, the expected final value can be observed.
- With larger loop counts, the program may produce incorrect and different results.
- Incrementing a shared counter requires multiple machine instructions.
- Because the operation is not atomic, thread interleavings can produce unexpected outcomes.

The Crux of the Problem

How can we build correct programs when many threads execute concurrently in the same memory space?

Persistence

- Main memory is volatile: data can disappear after power loss or a crash.
- Persistent storage keeps data for the long term.
- The file system manages persistent data on storage devices.
- Applications use system calls such as *open()*, *write()*, and *close()*.

```
1  #include <stdio.h>
2  #include <unistd.h>
3  #include <assert.h>
4  #include <fcntl.h>
5  #include <sys/types.h>
6
7  int main(int argc, char *argv[]) {
8      int fd = open("/tmp/file",
9                  O_WRONLY|O_CREAT|O_TRUNC,
10                 S_IRWXU);
11      assert(fd > -1);
12      int rc = write(fd, "hello world\n", 12);
13      assert(rc == 12);
14      close(fd);
15      return 0;
16  }
```

Figure 2.6: A Program That Does I/O (io.c)

Source: chapter page 10.

Reliable Persistent Storage

- File systems must determine where data is stored and maintain metadata about it.
- They issue I/O requests to underlying storage devices.
- For performance, writes may be delayed and grouped.
- Techniques such as **journaling** and **copy-on-write** help systems recover from failures.

The Crux of the Problem

How can the OS store data persistently with correctness, performance, and reliability?

Operating System Design Goals

- **Abstraction:** make the system convenient and manageable.
- **Performance:** minimize time and space overheads.
- **Protection and isolation:** prevent applications from harming one another or the OS.
- **Reliability:** keep the operating system running correctly.
- Additional goals may include energy efficiency, security, and mobility.

A Brief History of Operating Systems

1. **Early systems:** OS as libraries of common functions; batch processing.
2. **Protection:** system calls, user mode, kernel mode, and controlled privilege transitions.
3. **Multiprogramming:** several jobs in memory, rapid switching, better CPU utilization.
4. **UNIX:** simple, powerful abstractions and tools; strong influence on later systems.
5. **Modern era:** mature OS ideas returned to personal computers and later to mobile systems.

Three Major Themes

1. **Virtualization** of the CPU and memory.
2. **Concurrency** and correct coordination among simultaneous activities.
3. **Persistence** through devices and file systems.

- The OS provides abstractions and interfaces that make hardware easier to use.
- It manages resources while pursuing performance, protection, and reliability.

Virtualization

A Dialogue on Virtualization

A Dialogue on Virtualization

The dialogue introduces **virtualization** through the metaphor of a single physical peach being shared among several people, while each person has the illusion of having their own peach.

The same idea is applied to computing: a physical resource, such as a CPU, can be shared among multiple applications while each application behaves as if it had its own dedicated CPU.

The key idea is that the operating system creates this illusion by managing and sharing physical resources efficiently. In the case of the CPU, virtualization makes one physical processor appear as multiple virtual CPUs to running applications.

Thus, virtualization allows several programs to use the same hardware while giving each one the impression of exclusive access to the resource.

Processes

The Process Abstraction

- A **process** is a running program.
- A program stored on disk is passive:
 - instructions;
 - static data.
- The operating system transforms a program into an active process.
- Modern systems execute many processes seemingly at the same time.

Crux of the Problem

How can the OS provide the illusion of many CPUs when only a few physical CPUs are available?

Virtualizing the CPU

- The OS creates the illusion of many virtual CPUs.
- It runs one process, stops it, runs another, and so forth.
- This technique is called **time sharing**.
- CPU time is shared among concurrent processes.
- The main cost of sharing is reduced performance for each process.

Time Sharing

A resource is used by one entity for a period of time and then assigned to another.

Mechanisms and Policies

- CPU virtualization requires two important OS concepts:

Mechanisms

Low-level methods that implement functionality.

Example: stopping one process and starting another through a **context switch**.

Policies

Algorithms that determine which decision the OS should make.

Example: deciding which process should run next.

Mechanism = How?

Policy = Which?

What Defines a Process?

A process can be described through its **machine state**.

- **Address space**
 - program instructions;
 - program data.
- **CPU registers**
 - general-purpose registers;
 - program counter (PC);
 - stack pointer;
 - frame pointer.
- **I/O information**
 - open files and other resources.

A modern operating system provides interfaces to manage processes.

- **Create:** create a new process.
- **Destroy:** terminate a process.
- **Wait:** wait for a process to finish.
- **Miscellaneous Control:** suspend and resume processes.
- **Status:** obtain information about a process.

From Program to Process

- Programs initially reside on persistent storage.
- The OS loads:
 - program code;
 - static data.
- These elements are placed into the process address space.
- The process can then be prepared for execution.




Figure 4.1: Loading: From Program To Process

Source: chapter page 4.

Process Creation

Before starting execution, the OS performs several steps:

1. Load code and static data into memory.
2. Allocate the run-time **stack**.
3. Initialize arguments such as *argc* and *argv*.
4. Allocate initial space for the **heap**.
5. Initialize I/O resources.
6. Transfer control to the program entry point.

UNIX Processes

A process normally starts with three open file descriptors: standard input, standard output, and standard error.

Loading: Eager vs. Lazy

- Early or simple operating systems may load the entire program before execution.

Eager Loading

Code and data are loaded into memory all at once before execution.

Lazy Loading

Modern systems can load pieces of code and data only when they are needed during execution.

Lazy loading is related to memory virtualization mechanisms such as paging and swapping.

Process States

A simplified process model contains three states:

- **Running:** executing instructions on a CPU.
- **Ready:** able to run, but not currently selected.
- **Blocked:** waiting for an event, such as I/O completion.




Figure 4.2: **Process: State Transitions**

Source: chapter page 6

Process State Transitions

- **Ready → Running**
 - the process is **scheduled**.
- **Running → Ready**
 - the process is **descheduled**.
- **Running → Blocked**
 - the process initiates an operation such as I/O.
- **Blocked → Ready**
 - the event or I/O operation completes.

Example: CPU-Only Processes

- Consider two processes that only use the CPU.
- One process runs while the other remains ready.
- When the first process finishes, the second can run.

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process ₀ now done
5	–	Running	
6	–	Running	
7	–	Running	
8	–	Running	Process ₁ now done

Figure 4.3: **Tracing Process State: CPU Only**

Source: chapter page 6.

Example: CPU and I/O

- Process0 starts running.
- Process0 initiates I/O and becomes blocked.
- Process1 uses the CPU while Process0 waits.
- When I/O completes, Process0 becomes ready again.
- This improves CPU utilization.

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked, so Process ₁ runs
5	Blocked	Running	
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	–	
10	Running	–	Process ₀ now done

Figure 4.4: Tracing Process State: CPU and I/O

Source: chapter page 7.

Even a simple CPU/I/O example requires OS decisions.

- Which ready process should run?
- Should the OS switch processes when one becomes blocked?
- Should a process run immediately after its I/O completes?
- Should the currently running process continue?

Scheduler

The OS component responsible for deciding which process should use the CPU.

The OS must keep information about every process.

- process state;
- process identifier (PID);
- memory information;
- register context;
- open files;
- parent process;
- current directory.

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;
    int esp;
    int ebx;
    int ecx;
    int edx;
    int esi;
    int edi;
    int ebp;
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };

// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;           // Start of process memory
    uint sz;             // Size of process memory
    char *kstack;        // Bottom of kernel stack
                        // for this process
    enum proc_state state; // Process state
    int pid;             // Process ID
    struct proc *parent;  // Parent process
    void *chan;          // If !zero, sleeping on chan
    int killed;          // If !zero, has been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;    // Current directory
    struct context context; // Switch here to run process
    struct trapframe *tf; // Trap frame for the
                        // current interrupt
};
```

Figure 4.5: The xv6 Proc Structure

Process List and PCB

- The OS maintains a **process list** or **task list**.
- It tracks:
 - ready processes;
 - running processes;
 - blocked processes.
- Information about an individual process is often stored in a **Process Control Block (PCB)**.
- The saved register context allows the OS to stop and later resume a process.

Context Switch

The OS saves the state of one process and restores the state of another process.

Real operating systems may use more states than the simplified model.

- A process may have an initial state while it is being created.
- A process may remain temporarily after it finishes execution.

Zombie State

In UNIX-based systems, an exited process can remain in a zombie state until its parent obtains its return status, typically using `wait()`.

Chapter Summary

- A **process** is the OS abstraction of a running program.
- CPU virtualization allows many processes to share a small number of physical CPUs.
- A process is characterized by:
 - its address space;
 - CPU registers;
 - I/O state.
- Processes transition among states such as **running**, **ready**, and **blocked**.
- The OS tracks processes through structures such as the **process list** and **PCB**.
- Mechanisms and policies together enable CPU virtualization.

Process API

Key Question

What interfaces should the operating system provide for **process creation and control**?

- UNIX provides a distinctive process creation API.
- Three fundamental system calls:
 - *fork()* – create a new process;
 - *exec()* – execute a different program;
 - *wait()* – wait for a child process to finish.

The `fork()` System Call

- *fork()* creates a new process.
- The original process is the **parent**.
- The newly created process is the **child**.
- The child is an almost exact copy of the parent.
- Both continue execution after *fork()*.

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4
5  int main(int argc, char *argv[]) {
6      printf("hello (pid:%d)\n", (int) getpid());
7      int rc = fork();
8      if (rc < 0) {
9          // fork failed
10         fprintf(stderr, "fork failed\n");
11         exit(1);
12     } else if (rc == 0) {
13         // child (new process)
14         printf("child (pid:%d)\n", (int) getpid());
15     } else {
16         // parent goes down this path (main)
17         printf("parent of %d (pid:%d)\n",
18                rc, (int) getpid());
19     }
20     return 0;
21 }
```

Figure 5.1: Calling `fork()` (p1.c)

Return Value

- Parent: child PID.
- Child: 0.
- Failure: negative value.

Source: chapter page 2.

Parent and Child Execution

- Each process has its own:
 - address space;
 - registers;
 - program counter.
- Parent and child may execute in different orders.
- The scheduler determines which process runs first.
- Execution after *fork()* is therefore **non-deterministic**.

Process Identifier

Each process is identified by a unique **PID** (Process Identifier).

The `wait()` System Call

- A parent may need to wait for a child to finish.
- `wait()` blocks the parent.
- The parent becomes runnable again after the child terminates.
- This can make execution order predictable.

Execution Order

With `wait()`, the child completes before the parent continues past the system call.

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <sys/wait.h>
5
6  int main(int argc, char *argv[]) {
7      printf("hello (pid:%d)\n", (int) getpid());
8      int rc = fork();
9      if (rc < 0) {                // fork failed; exit
10         fprintf(stderr, "fork failed\n");
11         exit(1);
12     } else if (rc == 0) { // child (new process)
13         printf("child (pid:%d)\n", (int) getpid());
14     } else {                    // parent goes down this path
15         int rc_wait = wait(NULL);
16         printf("parent of %d (rc_wait:%d) (pid:%d)\n",
17                rc, rc_wait, (int) getpid());
18     }
19     return 0;
20 }
21
```

Figure 5.2: Calling `fork()` And `wait()` (p2.c)

Source: chapter page 3.

The `exec()` System Call

- `exec()` runs a different program.
- It replaces the current process image.
- Code and static data are replaced.
- Heap and stack are re-initialized.
- It does **not** create a new process.
- A successful `exec()` does not return.

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <string.h>
5 #include <sys/wait.h>
6
7 int main(int argc, char *argv[]) {
8     printf("hello (pid:%d)\n", (int) getpid());
9     int rc = fork();
10    if (rc < 0) {           // fork failed; exit
11        fprintf(stderr, "fork failed\n");
12        exit(1);
13    } else if (rc == 0) { // child (new process)
14        printf("child (pid:%d)\n", (int) getpid());
15        char *myargs[3];
16        myargs[0] = strdup("wc"); // program: "wc"
17        myargs[1] = strdup("p3.c"); // arg: input file
18        myargs[2] = NULL; // mark end of array
19        execvp(myargs[0], myargs); // runs word count
20        printf("this shouldn't print out");
21    } else {                // parent goes down this path
22        int rc_wait = wait(NULL);
23        printf("parent of %d (rc_wait:%d) (pid:%d)\n",
24              rc, rc_wait, (int) getpid());
25    }
26    return 0;
27 }
```

Figure 5.3: Calling `fork()`, `wait()`, And `exec()` (p3.c)

Source: chapter page 5.

Why Separate `fork()` and `exec()`?

- The separation is essential to the design of the UNIX shell.
- A shell typically:
 1. reads a command;
 2. calls `fork()`;
 3. modifies the child's environment if necessary;
 4. calls `exec()`;
 5. waits for completion with `wait()`.

Key Advantage

The shell can perform operations **between** `fork()` and `exec()`.

I/O Redirection

Example:

`wc p3.c > newfile.txt`

- The shell creates a child with *fork()*.
- The child closes standard output.
- It opens the destination file.
- The file becomes the new standard output.
- File descriptors remain open across *exec()*.

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <string.h>
5  #include <fcntl.h>
6  #include <sys/stat.h>
7  #include <sys/wait.h>
8
9  int main(int argc, char *argv[]) {
10     int rc = fork();
11     if (rc < 0) {
12         // fork failed
13         fprintf(stderr, "fork failed\n");
14         exit(1);
15     } else if (rc == 0) {
16         // child: redirect standard output to a file
17         close(STDOUT_FILENO);
18         open("./p4.output", O_CREAT|O_WRONLY|O_TRUNC,
19             S_IRWXU);
20         // now exec "wc"...
21         char *myargs[3];
22         myargs[0] = strdup("wc"); // program: wc
23         myargs[1] = strdup("p4.c"); // arg: file to count
24         myargs[2] = NULL; // mark end of array
25         execvp(myargs[0], myargs); // runs word count
26     } else {
27         // parent goes down this path (main)
28         int rc_wait = wait(NULL);
29     }
30     return 0;
31 }
```

Figure 5.4: All Of The Above With Redirection (p4.c)

Source: chapter page 8.

Pipes

- UNIX pipes use the *pipe()* system call.
- The output of one process is connected to an in-kernel pipe.
- The input of another process is connected to the same pipe.
- This allows commands to be combined into processing chains.

Example

```
grep -o foo file | wc -l
```

The output of *grep* becomes the input of *wc*.

Process Control and Signals

- UNIX provides additional interfaces for process control.
- *kill()* sends signals to processes.
- Signals can stop, continue, interrupt, or terminate processes.

Common Shell Shortcuts

- *Ctrl+C* → *SIGINT*
 - *Ctrl+Z* → *SIGTSTP*
 - *fg* can resume a stopped job.
- Processes can use *signal()* to react to signals.

- UNIX systems support multiple users.
- Users generally control only their own processes.
- The operating system protects processes belonging to different users.
- The **superuser** (*root*) has broader privileges.
- Root can control processes belonging to other users.

Superuser

Administrative privileges should be used carefully because they provide extensive control over the system.

- *ps* – displays running processes.
- *top* – displays processes and resource usage.
- *kill* – sends signals to processes.
- *killall* – another interface for sending signals.

Good Practice

Use UNIX *man* pages to study system calls, return values, options, and error conditions.

- *fork()* creates a new child process.
- *wait()* allows a parent to wait for a child.
- *exec()* replaces the current program with another.
- UNIX shells combine these mechanisms to execute commands.
- The separation of *fork()* and *exec()* enables redirection and pipes.
- Signals provide mechanisms for process control.
- Users and permissions determine which processes can be controlled.

Direct Execution

CPU Virtualization: The Main Challenge

Goal

The operating system must virtualize the CPU efficiently while retaining control over the system.

- The CPU is shared among multiple processes through **time sharing**.
- Two central challenges:
 - **Performance**: minimize virtualization overhead.
 - **Control**: prevent processes from taking over the system.
- Hardware and operating-system support are both required.

Basic Technique (Mechanism): Limited Direct Execution

- The basic idea is to run programs **directly on the CPU**.
- Direct execution provides high performance.
- Before execution, the OS:
 - creates a process entry;
 - allocates memory;
 - loads the program;
 - prepares the stack;
 - starts execution at the program entry point.

Problem

How can the OS run a program directly while still maintaining control over the CPU?

Direct Execution Without Limits

- The OS starts the program.
- The program executes directly on the processor.
- When it finishes, control returns to the OS.

Two problems emerge:

1. How to restrict privileged operations?
2. How to regain control of the CPU?

OS	Program
Create entry for process list Allocate memory for program Load program into memory Set up stack with argc/argv Clear registers Execute call main()	Run main() Execute return from main
Free memory of process Remove from process list	

Figure 6.1: Direct Execution Protocol (Without Limits)

Source: chapter page 2.

Problem #1: Restricted Operations

The Crux

A process must perform I/O and other restricted operations without gaining complete control over the system.

- Direct execution alone would allow applications to perform operations that should be controlled by the OS.
- The processor therefore supports different execution modes:
 - **User mode:** restricted execution.
 - **Kernel mode:** privileged execution.
- Privileged operations are performed by the operating system.

System Calls and Protected Control Transfer

- User programs request OS services through **system calls**.
- A system call uses a special **trap instruction**.
- The trap:
 1. saves the required processor state;
 2. switches execution to kernel mode;
 3. transfers control to a predefined OS handler.
- The OS performs the requested operation.
- A **return-from-trap** instruction restores execution in user mode.

Key Idea

System calls provide controlled access to privileged OS services.

Why System Calls Look Like Procedure Calls

- Calls such as *open()* and *read()* appear to be ordinary C procedure calls.
- The C library prepares:
 - system-call arguments;
 - the system-call number.
- Library code then executes the hardware-specific **trap instruction**.
- After returning from the kernel, the library processes the return value and returns control to the application.

The Trap Table

- User programs cannot choose arbitrary kernel addresses as trap destinations.
- During boot, the OS initializes a **trap table**.
- The table defines handlers for events such as:
 - system calls;
 - hardware interrupts;
 - exceptional conditions.
- Installing the trap table is a **privileged operation**.

Protection

User code requests a service through a system-call number, rather than specifying an arbitrary kernel address.

Limited Direct Execution Protocol

OS @ boot (kernel mode)	Hardware	
initialize trap table	remember address of... syscall handler	
OS @ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list Allocate memory for program Load program into memory Setup user stack with argv Fill kernel stack with reg/PC return-from-trap	restore regs (from kernel stack) move to user mode jump to main	Run main() ... Call system call trap into OS
Handle trap Do work of syscall return-from-trap	save regs (to kernel stack) move to kernel mode jump to trap handler	
	restore regs (from kernel stack) move to user mode jump to PC after trap	... return from main trap (via <code>exit()</code>)
Free memory of process Remove from process list		

Figure 6.2: Limited Direct Execution Protocol

Source: chapter page 5.

Problem #2: Switching Between Processes

The Crux

How can the operating system regain control of the CPU while a process is running?

- If a user process is executing, the OS itself is not executing.
- The OS therefore needs a mechanism to regain control.
- Two approaches are discussed:
 - cooperative;
 - non-cooperative.

- The OS waits for the process to voluntarily transfer control.
- Control can return to the OS through:
 - system calls;
 - an explicit *yield()* system call;
 - traps caused by illegal operations.

Limitation

A process running forever without making a system call could monopolize the CPU.

Non-Cooperative Approach: Timer Interrupt

- Hardware provides a programmable **timer interrupt**.
- The timer periodically interrupts the running process.
- When the interrupt occurs:
 1. the running process is halted;
 2. its state is saved;
 3. the CPU switches to kernel mode;
 4. the OS interrupt handler executes.
- The OS has regained control of the CPU.

Key Mechanism

The timer interrupt prevents a process from monopolizing the CPU.

- After regaining control, the OS scheduler decides whether:
 - to continue executing the current process; or
 - to switch to another process.
- A **context switch** saves the state of one process and restores the state of another.
- Relevant state includes:
 - general-purpose registers;
 - program counter;
 - kernel stack pointer.

Timer Interrupt and Context Switch

- Process A is executing.
- A timer interrupt occurs.
- Hardware saves A's user state.
- The OS enters the interrupt handler.
- The OS switches from A to B.
- B resumes execution.




Figure 6.3: Limited Direct Execution Protocol (Timer Interrupt)

Source: chapter page 11.

Two Types of State Saving

1. Hardware save/restore

- occurs when the timer interrupt happens;
- saves the user registers;
- uses the process kernel stack.

2. Software save/restore

- occurs during the context switch;
- performed explicitly by the OS;
- saves and restores kernel context.

Example: xv6 Context Switch

- Save registers of the old process.
- Restore registers of the new process.
- Switch the kernel stack.
- Return into the new execution context.

```
1 # void swtch(struct context *old, struct context *new);
2 #
3 # Save current register context in old
4 # and then load register context from new.
5 .globl swtch
6 swtch:
7     # Save old registers
8     movl 4(%esp), %eax # put old ptr into eax
9     popl 0(%eax)      # save the old IP
10    movl %esp, 4(%eax) # and stack
11    movl %ebx, 8(%eax) # and other registers
12    movl %ecx, 12(%eax)
13    movl %edx, 16(%eax)
14    movl %esi, 20(%eax)
15    movl %edi, 24(%eax)
16    movl %ebp, 28(%eax)
17
18    # Load new registers
19    movl 4(%esp), %eax # put new ptr into eax
20    movl 28(%eax), %ebp # restore other registers
21    movl 24(%eax), %edi
22    movl 20(%eax), %esi
23    movl 16(%eax), %edx
24    movl 12(%eax), %ecx
25    movl 8(%eax), %ebx
26    movl 4(%eax), %esp # stack is switched here
27    pushl 0(%eax)      # return addr put in place
28    ret                # finally return into new ctxt
```

Figure 6.4: The xv6 Context Switch Code

What About Concurrency?

- What happens if a timer interrupt occurs during a system call?
- What happens if another interrupt arrives while one is being handled?

The chapter briefly introduces two approaches:

- temporarily disabling interrupts during interrupt processing;
- using locking mechanisms to protect kernel data structures.

Next Step

Detailed treatment of these problems belongs to the study of **concurrency**.

Limited Direct Execution: Summary

- Programs execute directly on the CPU for high performance.
- **User mode** restricts application privileges.
- **Kernel mode** allows privileged operations.
- **System calls** use traps to enter the kernel safely.
- A **trap table** determines valid kernel entry points.
- A **timer interrupt** allows the OS to regain CPU control.
- A **context switch** changes the running process.

Central Idea

Limited direct execution combines efficient execution with operating-system control.

Performance Measurement

The chapter suggests measuring two fundamental operating-system costs:

- **System-call cost**
 - repeatedly execute a simple system call;
 - measure total execution time;
 - estimate the average cost per call.
- **Context-switch cost**
 - use two processes communicating through UNIX pipes;
 - force repeated switches between the processes;
 - measure the communication cycle.

Tools Mentioned in the Chapter

gettimeofday(), *rdtsc*, *sched_setaffinity()*, and *lmbench*.

Direct Execution: Summary

- **Limited Direct Execution** combines performance and OS control.
- **User/kernel modes** restrict privileged operations.
- **System calls** use traps to safely enter the kernel.
- **Timer interrupts** allow the OS to regain CPU control.
- **Context switches** enable execution to move between processes.

Next Step

With these mechanisms in place, the remaining question is:

Which process should run at a given time?

This is the responsibility of the **CPU scheduler**.

CPU Scheduling

From Mechanisms to Scheduling

- We now understand the low-level mechanisms for running processes:
 - timer interrupts;
 - context switching;
 - switching between processes.
- The next question concerns **policy**:
 - Which process should run?
 - When should it run?
 - For how long?

The Crux

How should we develop a basic framework for thinking about **scheduling policies**?

To develop basic scheduling policies, we initially assume:

1. Each job runs for the same amount of time.
2. All jobs arrive at the same time.
3. Once started, each job runs to completion.
4. Jobs use only the CPU and perform no I/O.
5. The run-time of each job is known.

These assumptions are intentionally unrealistic and will be **relaxed progressively**.

The main performance metric initially considered is **turnaround time**:

$$T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$$

With all jobs arriving at $t = 0$:

$$T_{\text{turnaround}} = T_{\text{completion}}$$

Performance vs. Fairness

Scheduling policies may improve performance at the cost of fairness, or improve fairness at the cost of performance.

First In, First Out (FIFO)

- Also called **First Come, First Served (FCFS)**.
- Simple and easy to implement.
- Jobs execute in arrival order.




Figure 7.1: **FIFO Simple Example**

Source: chapter page 3.

Example:

$$A = B = C = 10$$

$$\bar{T}_{turnaround} = \frac{10 + 20 + 30}{3} = 20$$

The Convoy Effect

Relax assumption 1: jobs may have **different lengths**.

Example:

$$A = 100, \quad B = C = 10$$

Under FIFO:

$$\bar{T}_{\text{turnaround}} = \frac{100 + 110 + 120}{3} = 110$$




Figure 7.2: **Why FIFO Is Not That Great**

Source: chapter page 3.

Convoy Effect

Short jobs wait behind a long-running job.

Shortest Job First (SJF)

Shortest Job First (SJF)
executes jobs in increasing
order of run time.

For:

$$A = 100, \quad B = C = 10$$

SJF executes:

$$B \rightarrow C \rightarrow A$$

and obtains:

$$\bar{T}_{\text{turnaround}} = \frac{10 + 20 + 120}{3} = 50$$




Figure 7.3: **SJF Simple Example**

Source: chapter page 4.

SJF and Late Arrivals

Now relax assumption 2: jobs may arrive at **different times**.

Example:

- A arrives at $t = 0$, requiring 100 s.
- B and C arrive at $t = 10$, requiring 10 s each.

Because SJF is non-preemptive, B and C must still wait for A.




Figure 7.4: **SJF With Late Arrivals From B and C**

Source: chapter page 5.

$$\bar{T}_{turnaround} = 103.33$$

Shortest Time-to-Completion First (STCF)

Relax assumption 3: a running job may be **preempted**.

STCF is the preemptive version of SJF.

- When a new job arrives, compare the remaining execution times.
- Run the job with the least time remaining.
- A running process may therefore be preempted.

For the previous example:

$$\bar{T}_{turnaround} = 50$$




Figure 7.5: **STCF Simple Example**

Source: chapter page 6.

A New Metric: Response Time

Interactive systems require another metric: **response time**.

$$T_{\text{response}} = T_{\text{firstrun}} - T_{\text{arrival}}$$

- Turnaround time focuses on **completion**.
- Response time focuses on how quickly a job **starts receiving CPU time**.
- SJF/STCF may provide excellent turnaround time but poor interactive response.

Round Robin (RR)

Round Robin introduces **time slicing**.

- Each job runs for a fixed **time slice** or **scheduling quantum**.
- After the quantum expires, the scheduler switches to the next job in the run queue.
- The process repeats until all jobs finish.

Time-Slice Trade-off

A shorter time slice improves response time, but increases the relative cost of context switching.

- Example: 10 ms quantum + 1 ms context switch $\Rightarrow$ roughly 10% switching overhead.
- Increasing the quantum amortizes this fixed cost.

Turnaround Time vs. Response Time

The scheduling policies studied so far optimize different goals.

Policy	Turnaround	Response
FIFO	workload-dependent	poor with convoy effect
SJF	very good	potentially poor
STCF	very good	potentially poor
RR	poor	very good

Fundamental Trade-off

Policies that divide CPU time fairly among active processes tend to improve response time, but usually increase turnaround time.

Relax assumption 4: processes may perform I/O.

- A process issuing I/O becomes blocked.
- The scheduler can run another process while the I/O proceeds.
- When the I/O completes, the blocked process becomes ready again.

Example:

- Job A: five 10-ms CPU bursts separated by I/O.
- Job B: one continuous 50-ms CPU burst.

Overlapping CPU and I/O

Poor Resource Use




Figure 7.8: **Poor Use Of Resources**

Source: chapter page 10.

CPU/I/O Overlap




Figure 7.9: **Overlap Allows Better Use Of Resources**

Source: chapter page 10.

Key Idea

Treating each CPU burst as a separate job enables CPU execution to overlap with another process's I/O, improving system utilization.

No More Oracle

One assumption still remains:

The scheduler knows the run-time of each job.

In a general-purpose operating system, this information is normally **not known in advance**.

The remaining challenge is therefore:

- obtain SJF/STCF-like turnaround behavior;
- preserve good response time;
- operate without knowing future job lengths.

Next Step

Use the **recent past to predict the future**: the Multi-Level Feedback Queue (MLFQ).

CPU Scheduling: Summary

- Scheduling policy determines **which process runs**.
- **FIFO** is simple but may suffer from the convoy effect.
- **SJF/STCF** optimize turnaround time.
- **Round Robin** improves response time through time slicing.
- Response time and turnaround time represent an inherent scheduling trade-off.
- CPU and I/O can be overlapped to improve utilization.
- Real systems must schedule without knowing future job lengths.

Next: Multi-Level Feedback Queue (MLFQ)

References



Arpaci-Dusseau, Remzi H. and Andrea C. Arpaci-Dusseau.
Operating Systems: Three Easy Pieces. 1.10. Arpaci-Dusseau Books,
Nov. 2023. URL: <https://pages.cs.wisc.edu/~remzi/OSTEP/>.

? & | !